

Hands on With SQL

Obectives

- Define and Understand DDL and DML
- Create a Simple Schema and Seed data in .sql file
- Understand how SQL handles CRUD
- Dump, then Delete you DB then restore it from the dump file
- SQL query practice with JOINS
- Understand and use 3 Aggregate functions

DDL Data Definition Language

(things like CREATE TABLE, DROP TABLE, ALTER, RENAME, TRUNCATE (empty for re-use) and defining data types in SQL fields)

DML Data Manipulation Language

(used in queries, especially CRUD operations)

How does SQL Handle CRUD?

C

R

U

D

How does SQL Handle CRUD?

- **C** INSERT INTO
- **R** SELECT
- **U** UPDATE, SET
- **D** DELETE FROM

EXAMPLES

INSERT (CREATE)

```
INSERT INTO CUSTOMERS (ID,NAME,AGE,ADDRESS,SALARY)
VALUES (1, 'Ramesh', 32, 'Ahmedabad', 2000.00 );
```

```
INSERT INTO CUSTOMERS (ID,NAME,AGE,ADDRESS,SALARY)
VALUES (2, 'Khilan', 25, 'Delhi', 1500.00 );
```

```
INSERT INTO CUSTOMERS (ID,NAME,AGE,ADDRESS,SALARY)
VALUES (3, 'kaushik', 23, 'Kota', 2000.00 );
```

```
INSERT INTO CUSTOMERS (ID,NAME,AGE,ADDRESS,SALARY)
VALUES (4, 'Chaitali', 25, 'Mumbai', 6500.00 );
```

```
INSERT INTO CUSTOMERS (ID,NAME,AGE,ADDRESS,SALARY)
VALUES (5, 'Hardik', 27, 'Bhopal', 8500.00 );
```

```
INSERT INTO CUSTOMERS (ID,NAME,AGE,ADDRESS,SALARY)
VALUES (6, 'Komal', 22, 'MP', 4500.00 );
```

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
6	Komal	22	MP	4500.00

SELECT (READ)

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
6	Komal	22	MP	4500.00
7	Muffy	24	Indore	10000.00

SELECT ID, NAME, SALARY FROM CUSTOMERS

ID	NAME	SALARY
1	Ramesh	2000.00
2	Khilan	1500.00
3	kaushik	2000.00
4	Chaitali	6500.00
5	Hardik	8500.00
6	Komal	4500.00
7	Muffy	10000.00

UPDATE (UPDATE)

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
6	Komal	22	MP	4500.00
7	Muffy	24	Indore	10000.00

```
UPDATE CUSTOMERS
SET ADDRESS = 'Pune'
WHERE ID = 6;`
```

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
6	Komal	22	Pune	4500.00
7	Muffy	24	Indore	10000.00

DELETE (DELETE)

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
6	Komal	22	MP	4500.00
7	Muffy	24	Indore	10000.00

DELETE FROM CUSTOMERS

WHERE ID = 6;

Now, CUSTOMERS table would have the following records:

ID	NAME	AGE	ADDRESS	SALARY
1	Ramesh	32	Ahmedabad	2000.00
2	Khilan	25	Delhi	1500.00
3	kaushik	23	Kota	2000.00
4	Chaitali	25	Mumbai	6500.00
5	Hardik	27	Bhopal	8500.00
7	Muffy	24	Indore	10000.00

I have developed a
Zen-like approach to
the operating systems
that people use:
'When you're ready,
the right operating
system will appear in
your life.

Guy Kawasaki

Create your schema using .sql files

- `createdb weatherapp`
- `mkdir SQL_FUN`
- `//cd into SQL_FUN`
- `touch db/schema.sql`
- `touch db/seed.sql //seed data`

...in your schema

```
DROP TABLE IF EXISTS weather;
```

```
DROP TABLE IF EXISTS cities;
```

```
CREATE TABLE cities (  
    id serial primary key,  
    city varchar(80),  
    location point  
);
```

```
CREATE TABLE weather (  
    id serial primary key,  
    city_id int references cities(id) on delete cascade,  
    temp_hi int CHECK (temp_hi > temp_lo),  
    temp_lo int,  
    prcp real,  
    date date  
);
```

...in your seed file

```
INSERT INTO
```

```
  cities
```

```
VALUES
```

```
  ( default, 'Boulder', point(2,5) ),  
  ( default, 'Denver', point(7,2) ),  
  ( default, 'Brooklyn', point(9,1) );
```

```
INSERT INTO
```

```
  weather
```

```
VALUES
```

```
  ( default, (SELECT id FROM cities WHERE city = 'Boulder'), 75, 42, 210000, now() ),  
  ( default, (SELECT id FROM cities WHERE city = 'Denver'), 65, 55, 300030, now() ),  
  ( default, (SELECT id FROM cities WHERE city = 'Brooklyn'), 55, 39, 120000, now() ),  
  ( default, (SELECT id FROM cities WHERE city = 'Boulder'), 71, 55, 103000, ( now() - interval '1' day ) ),  
  ( default, (SELECT id FROM cities WHERE city = 'Denver'), 74, 51, 300040, ( now() - interval '1' day ) ),  
  ( default, (SELECT id FROM cities WHERE city = 'Brooklyn'), 72, 66, 203000, ( now() - interval '1' day ) ),  
  ( default, (SELECT id FROM cities WHERE city = 'Boulder'), 81, 65, 104000, ( now() - interval '2' day ) ),  
  ( default, (SELECT id FROM cities WHERE city = 'Denver'), 64, 55, 300300, ( now() - interval '2' day ) ),  
  ( default, (SELECT id FROM cities WHERE city = 'Brooklyn'), 42, 36, 202000, ( now() - interval '2' day ) );
```

. . . from the command line

- `psql weatherapp > db/schema.sql`
- `psql weatherapp > db/seed.sql`
- `psql weatherapp`
- `\dt // JOY!!! you have tables`
- `\d cities // check your schema and foreign keys`

JOINS

Joins are where SQL gets magical.

Joins are how you utilize the relational part of the database. d

Without joins, your database is essentially useless. :(

LET'S JOIN SOME RELATIONAL
DATA together

There are many joins, but INNER JOIN
is your friend

Anatomy of a SQL JOIN

```
SELECT <select_list>  
FROM Table_A  
INNER JOIN Table_B  
ON A.Key = B.Key
```

SEMANTIC aside to make your joins easier...

- call the PK in any table 'id'
- when you reference that in another table, name it 'tableName'_id

EXAMPLE

users		posts		comments	
-----		-----		-----	
id	PK	id	PK	id	PK
name		user_id	FK	user_id	FK
age		body		post_id	FK

This solves the 'ambiguous' error you may have run into

BACK TO JOINS

1. LETS JOIN cities and Weather (we do)
(Ya'll do)
2. Write a Query to display the date and the city name where the high was over 70
3. Write a Query to display the date and the city name where the temp range was less than 8 degrees

Dumping and Restoring

DUMP

```
pg_dump dbname > outfile.sql
```

RESTORE

```
psql dbname < infile
```

Aggregate functions

(simply, they aggregate data...things like SUM, MIN, MAX, COUNT, AVG)

SQL can also compute values on the fly

ie)

column 1 :population

column 2 :GDP

column 3 (computed): population/GDP AS perCapita GDP

SQL CHEATSHEET

Basic Queries

- filter your columns
SELECT col1, col2, col3, ... **FROM** table1
- filter the rows
WHERE col4 = 1 **AND** col5 = 2
- aggregate the data
GROUP by ...
- limit aggregated data
HAVING count(*) > 1
- order of the results
ORDER BY col2

Useful keywords for **SELECTS**:

- DISTINCT** - return unique results
- BETWEEN** a **AND** b - limit the range, the values can be numbers, text, or dates
- LIKE** - pattern search within the column text
- IN** (a, b, c) - check if the value is contained among given.

Data Modification


- update specific data with the **WHERE** clause
UPDATE table1 **SET** col1 = 1 **WHERE** col2 = 2
- insert values manually
INSERT INTO table1 (ID, FIRST_NAME, LAST_NAME)
VALUES (1, 'Rebel', 'Labs');
- or by using the results of a query
INSERT INTO table1 (ID, FIRST_NAME, LAST_NAME)
SELECT id, last_name, first_name **FROM** table2

Views

A **VIEW** is a virtual table, which is a result of a query. They can be used to create virtual tables of complex queries.

```
CREATE VIEW view1 AS
SELECT col1, col2
FROM table1
WHERE ...
```


The Joy of JOINS



LEFT OUTER JOIN - all rows from table A, even if they do not exist in table B



INNER JOIN - fetch the results that exist in both tables



RIGHT OUTER JOIN - all rows from table B, even if they do not exist in table A

Updates on JOINed Queries

You can use **JOINS** in your **UPDATES**

```
UPDATE t1 SET a = 1
FROM table1 t1 JOIN table2 t2 ON t1.id = t2.t1_id
WHERE t1.col1 = 0 AND t2.col2 IS NULL;
```

NB! Use database specific syntax, it might be faster!

Semi JOINS

You can use subqueries instead of **JOINS**:

```
SELECT col1, col2 FROM table1 WHERE id IN
(SELECT t1_id FROM table2 WHERE date >
CURRENT_TIMESTAMP)
```

Indexes

If you query by a column, index it!
CREATE INDEX index1 **ON** table1 (col1)

Don't forget:

Avoid overlapping indexes

Avoid indexing on too many columns

Indexes can speed up **DELETE** and **UPDATE** operations

Useful Utility Functions

- convert strings to dates:
TO_DATE (Oracle, PostgreSQL), **STR_TO_DATE** (MySQL)
- return the first non-NULL argument:
COALESCE (col1, col2, "default value")
- return current time:
CURRENT_TIMESTAMP
- compute set operations on two result sets
SELECT col1, col2 **FROM** table1
UNION / EXCEPT / INTERSECT
SELECT col3, col4 **FROM** table2;

Union - returns data from both queries
Except - rows from the first query that are not present in the second query
Intersect - rows that are returned from both queries

Reporting

Use aggregation functions

- COUNT** - return the number of rows
- SUM** - cumulate the values
- AVG** - return the average for the group
- MIN / MAX** - smallest / largest value

More practice

- JOINS on SQL Zoo!
- Query Stack Overflow

[https://data.stackexchange.com/stackoverflow/
query/new](https://data.stackexchange.com/stackoverflow/query/new)

Coming soon . . .

- new information to make all of this easier and nearly irrelevant!!!